

Socket programming

From theory to implementation


What you will know after this tutorial

- How computers communicate on the network
 - What are sockets? Why would you use them?
- How to *practically* use sockets to enable the communication
 - How do you do this in C?

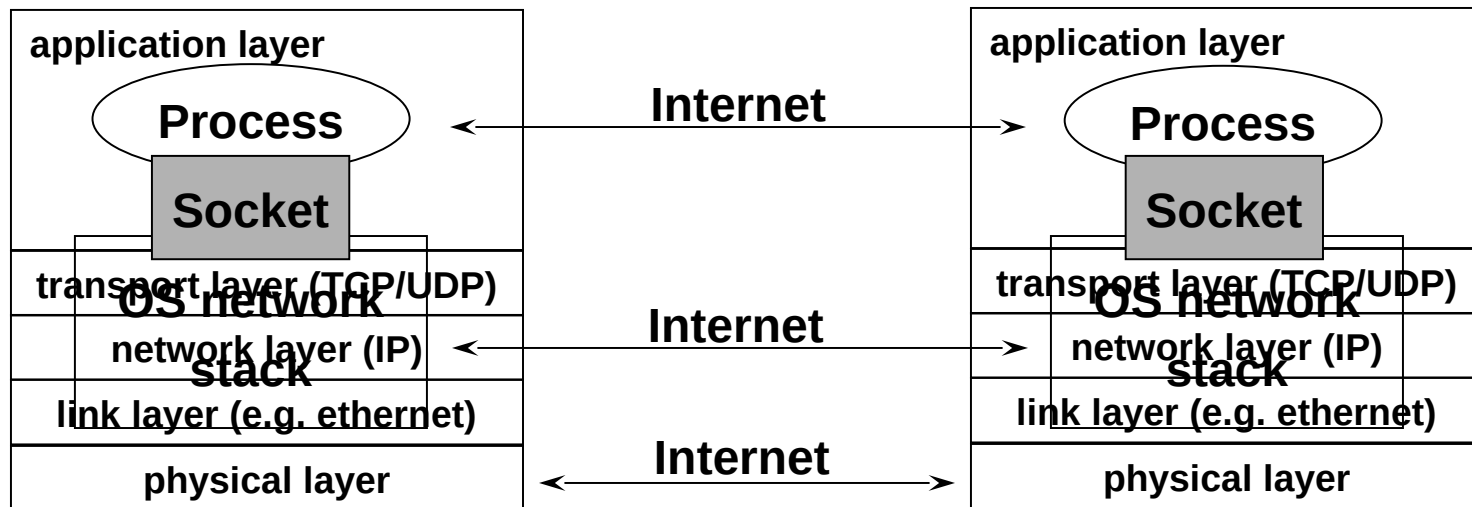
Socket programming is simple

Client/sever model

- Client asks (*request*) – server provides (*response*)
- Typically: single server - multiple clients
- The server does not need to know *anything* about the client
 - even that it exists
- The client should always know *something* about the server
 - at least where it is located

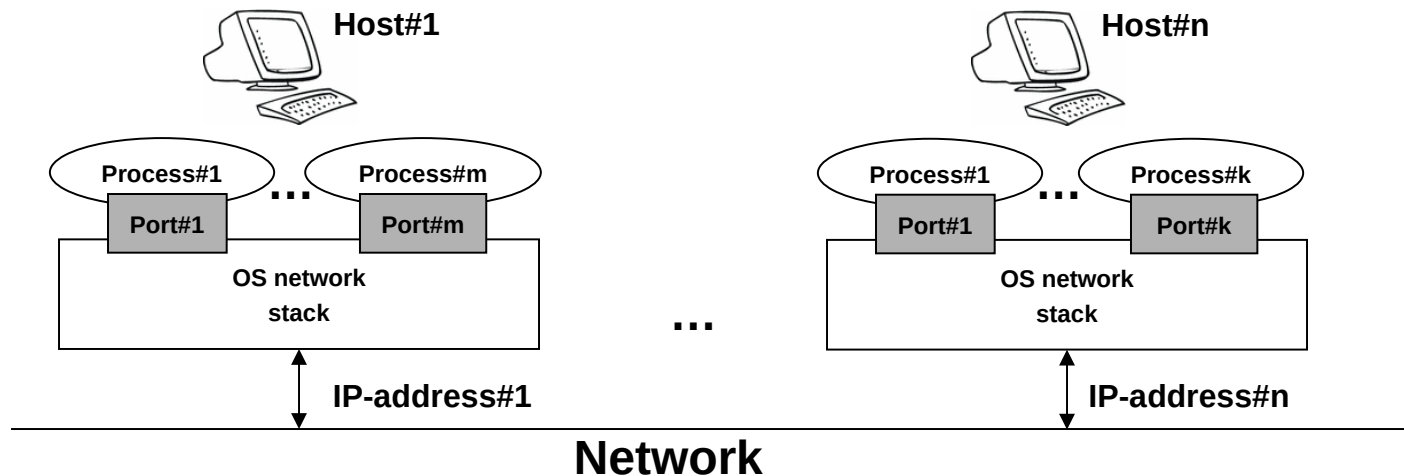


Sockets as means for inter-process communication (IPC)



Addressing server

- Address the machine on the network
 - By IP address (127.0.0.1 – localhost)
- Address the process
 - By the “port”-number
- The pair of *IP-address* + *port* – makes up a “*socket-address*”

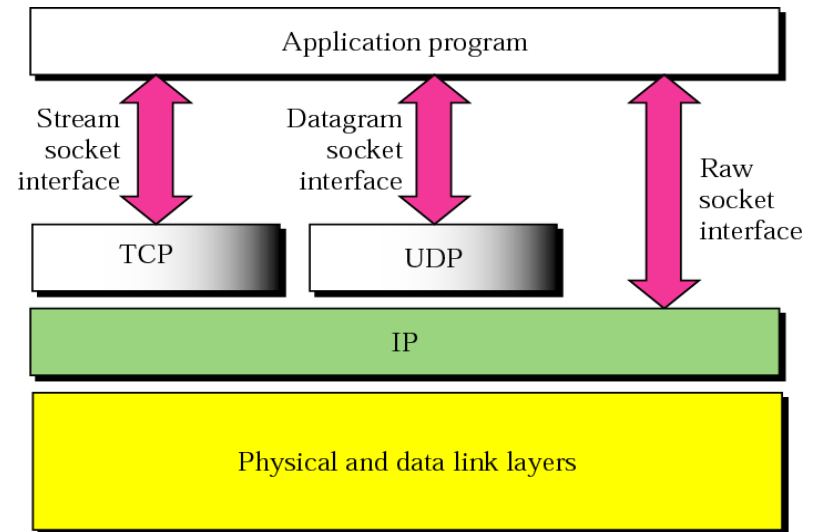


Usage of port-numbers

- Standard applications use predefined port-numbers
 - 21 - ftp
 - 23 - telnet
 - 80 - http
 - 110 - pop3 (email)
 - ...
- Other applications should choose between 1024 and 65535
 - 4662 – eMule
 - ...

Socket types

- Datagram socket – using UDP
 - Not sequenced
 - Not reliable
 - Not unduplicated
 - Connectionless
 - Border preserving
- Stream socket – using TCP
 - Sequenced
 - Reliable
 - Unduplicated
 - Connection-oriented
 - Not border preserving
- Raw and others (extracurricular)



Let's summarize

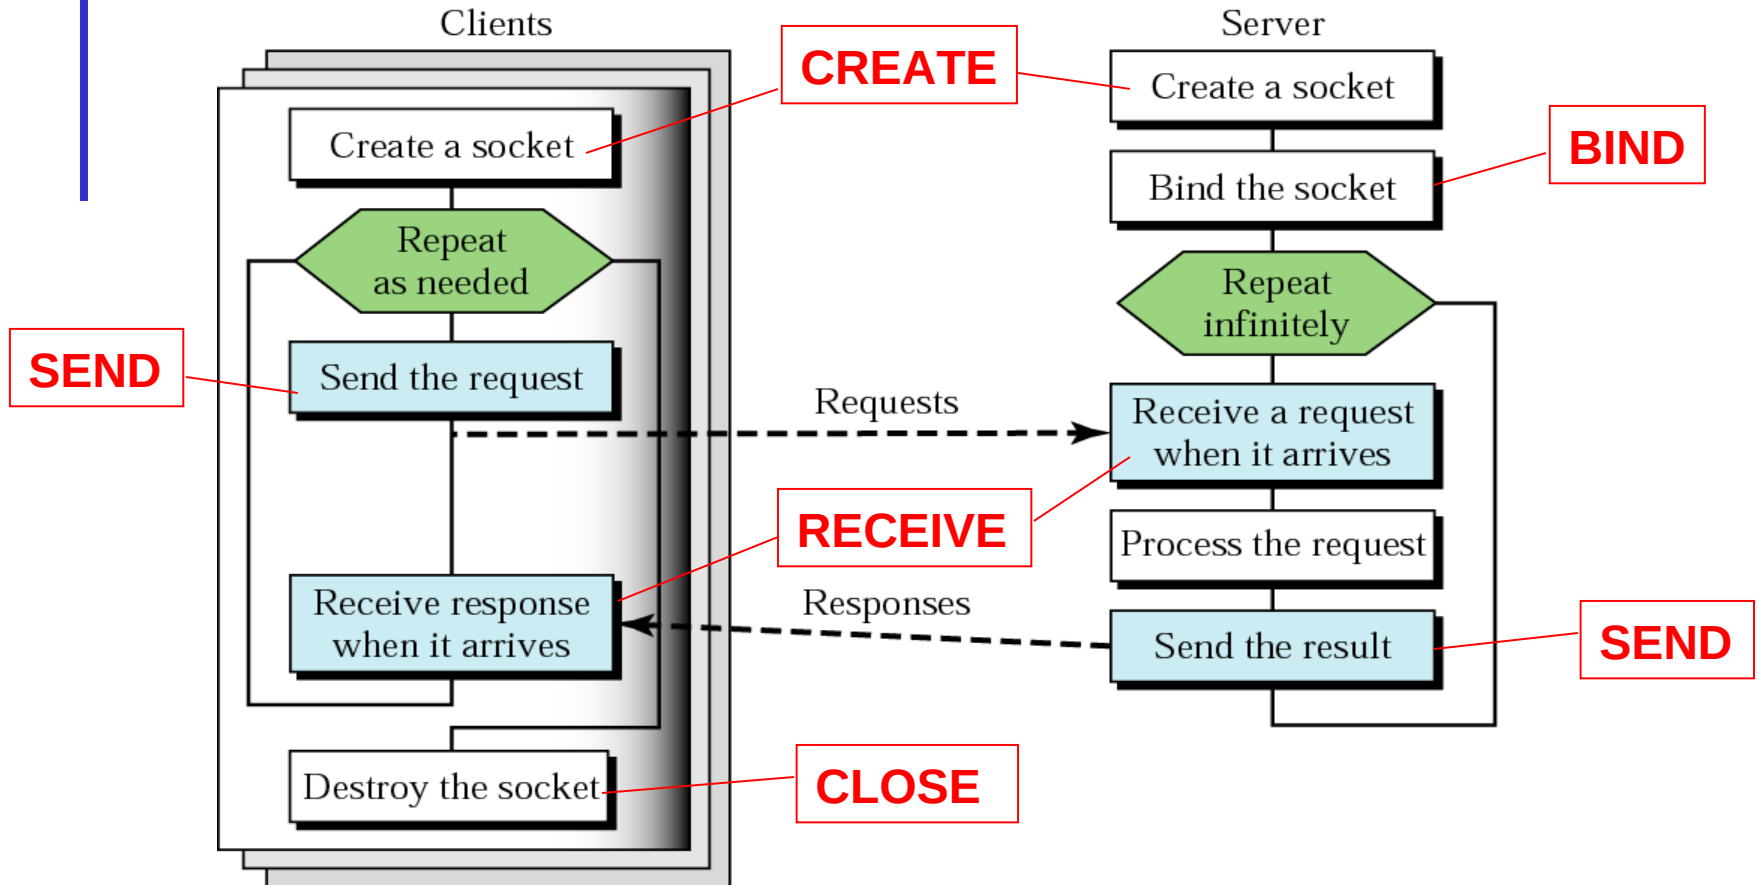
- What is a *socket*?
 - A *socket* is a communication end-point between two processes (which are normally located on different machines)
- How to address a socket?
 - By a *socket address*: IP-address + port number
- Where are sockets used?
 - Any network application: ftp-applications, web-browsers, web-servers, telnet servers/terminals, email-server/clients, P2P-applications (Kazaa, eMule etc), chat-clients (ICQ, Messenger) and many more
- Which programming language do I need to program sockets?
 - Almost any language will do

Socket primitives

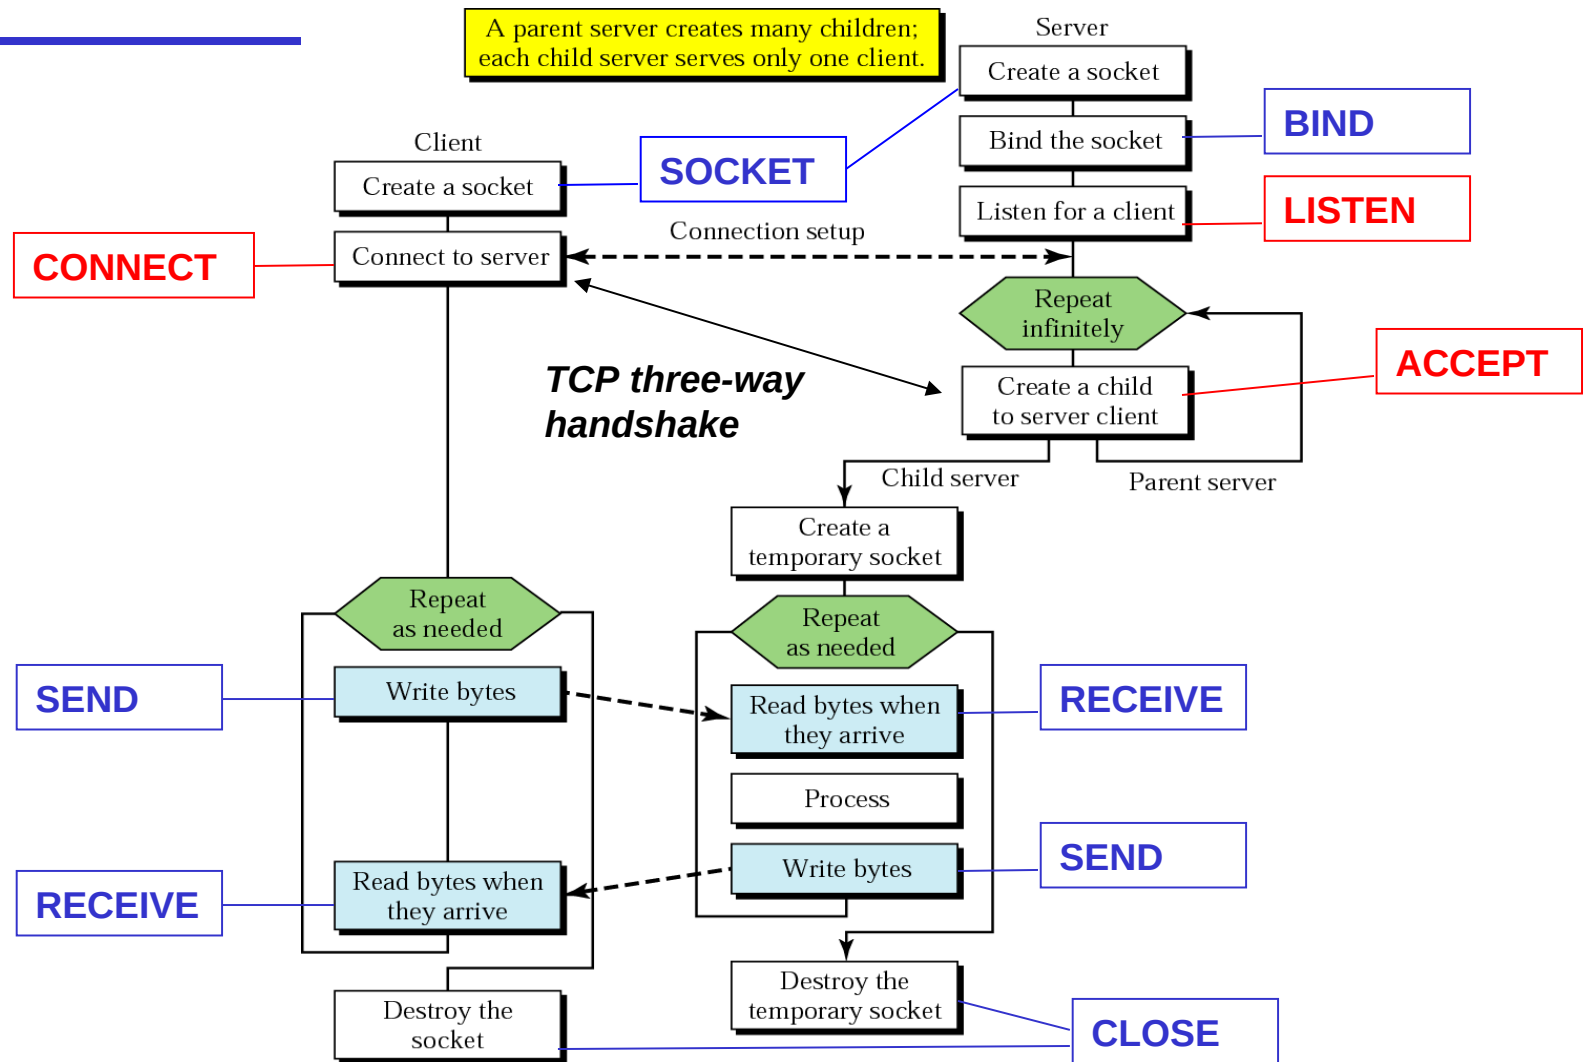
- SOCKET create a new socket
- BIND attach a local address to a socket
- LISTEN announce a willingness to accept connections
- ACCEPT block the caller process until a connection attempt arrives
- CONNECT connection actively attempt to establish a connection
- SEND send some data over the connection
- RECEIVE connection receive some data from the connection
- CLOSE release the connection (the port)

Client+server: connectionless

Each server serves many clients but handles one request at a time.



Client+server: connection-oriented



Primitives in C

- **SOCKET**: `int socket(int domain, int type, int protocol);`
 - `domain := AF_INET`
 - `type := (SOCK_DGRAM or SOCK_STREAM)`
 - `protocol := 0`
 - `returned`: socket descriptor (`sockfd`)
- **BIND**: `int bind(int sockfd, struct sockaddr *my_addr, int addrlen);`
 - `sockfd` - socket descriptor (returned from `socket()`)
 - `my_addr`: socket address
 - `addrlen := sizeof(struct sockaddr)`

“-1” returned?

- a problem!

```
my_addr.sin_family = AF_INET;  
my_addr.sin_port = 0; // choose an unused port at random  
my_addr.sin_addr.s_addr = INADDR_ANY; // use my IP address  
memset(&(my_addr.sin_zero), '\0', 8); // zero the rest of the struct
```

Primitives in C (continued)

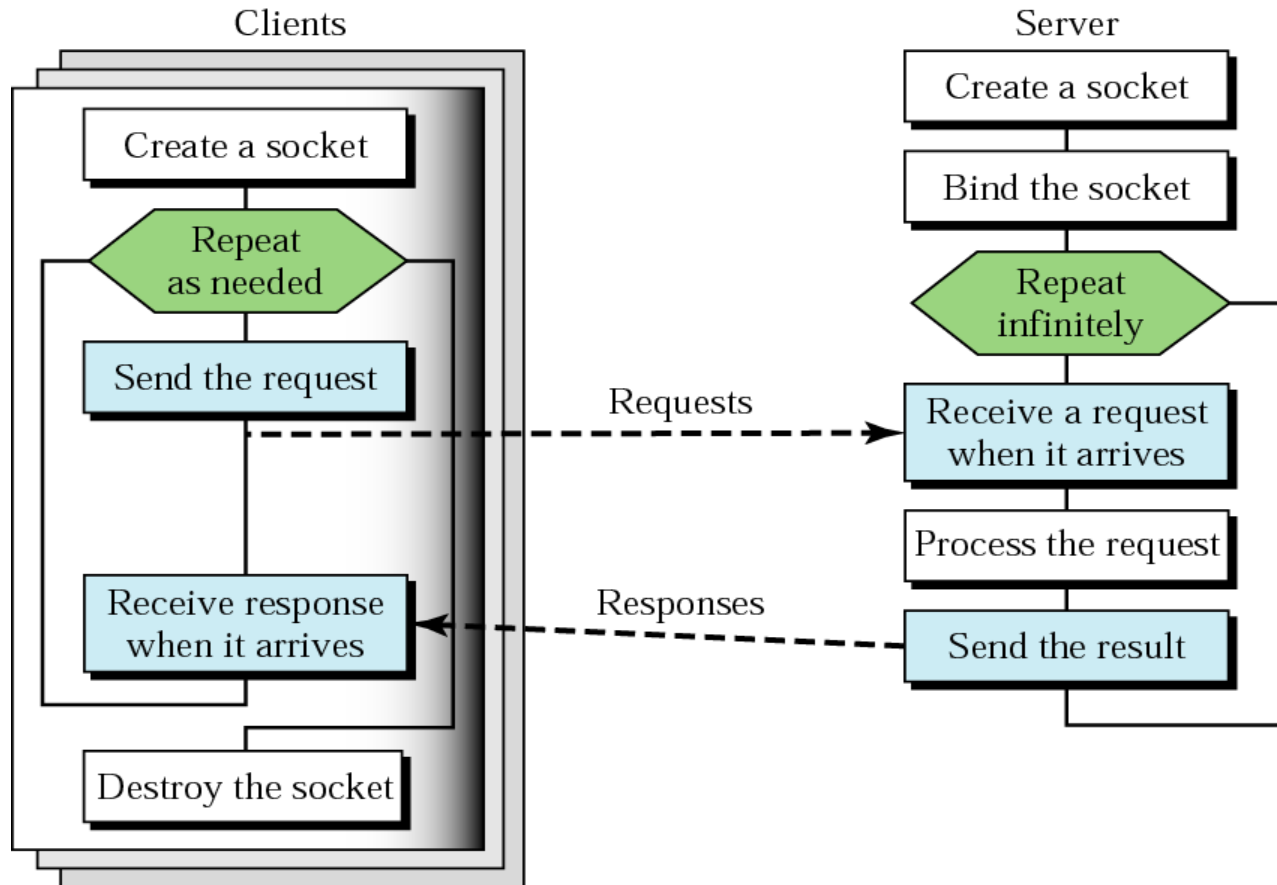
- **LISTEN**: `int listen(int sockfd, int backlog);`
 - *backlog*: how many connections we want to queue
- **ACCEPT**: `int accept(int sockfd, void *addr, int *addrlen);`
 - *addr*: here the socket-address of the caller will be written (use *struct sockaddr*)
 - *returned*: a new socket descriptor (for the temporal socket)
- **CONNECT**: `int connect(int sockfd, struct sockaddr *serv_addr, int addrlen);`
 - parameters are same as for `bind()`
- **SEND**: `int send(int sockfd, const void *msg, int len, int flags);`
 - *msg*: message you want to send
 - *len*: length of the message
 - *flags* := 0
 - *returned*: the number of bytes actually sent
- **RECEIVE**: `int recv(int sockfd, void *buf, int len, unsigned int flags);`
 - *buf*: buffer to receive the message
 - *len*: length of the buffer (“don’t give me more!”)
 - *flags* := 0
 - *returned*: the number of bytes received

Primitives in C (continued)


- **SEND** (DGRAM-style): `int sendto(int sockfd, const void *msg, int len, int flags, const struct sockaddr *to, int tolen);`
 - *msg*: message you want to send
 - *len*: length of the message
 - *flags* := 0
 - *to*: socket address of the remote process
 - *tolen*: = sizeof(struct sockaddr)
 - *returned*: the number of bytes actually sent
- **RECEIVE** (DGRAM-style): `int recvfrom(int sockfd, void *buf, int len, unsigned int flags, struct sockaddr *from, int *fromlen);`
 - *buf*: buffer to receive the message
 - *len*: length of the buffer (“don’t give me more!”)
 - *from*: socket address of the process that sent the data
 - *fromlen*:= sizeof(struct sockaddr)
 - *flags* := 0
 - *returned*: the number of bytes received
- **CLOSE**: `close (socketfd);`

Client+server: connectionless (repeated)

Each server serves many clients but handles one request at a time.



Example application: “Echo” (*in ANSI-C*)



- Client sends a message to the server
- Server echos this message back to the client
- *We will use datagram sockets*

EchoClient.c – #include's and #define's

```
#include <stdio.h>    /* for printf() and fprintf() */
#include <sys/socket.h> /* for socket(), connect(), sendto(), and recvfrom() */
#include <arpa/inet.h> /* for sockaddr_in and inet_addr() */
#include <stdlib.h>    /* for atoi() and exit() */
#include <string.h>    /* for memset() */
#include <unistd.h>    /* for close() */

#define ECHOMAX 255    /* Longest string to echo */

void DieWithError(char *errorMessage); /* External error handling function */
```

EchoClient.c – variable declarations

```
int main(int argc, char *argv[])
{
    int sock;                /* Socket descriptor */
    struct sockaddr_in echoServAddr; /* Echo server address */
    struct sockaddr_in fromAddr;    /* Source address of echo */
    unsigned short echoServPort;   /* Echo server port */
    unsigned int fromSize;         /* In-out of address size for recvfrom() */
    char *servIP;                 /* IP address of server */
    char *echoString;             /* String to send to echo server */
    char echoBuffer[ECHOMAX+1];   /* Buffer for receiving echoed string */
    int echoStringLen;            /* Length of string to echo */
    int respStringLength;         /* Length of received response */

    if ((argc < 3) || (argc > 4)) /* Test for correct number of arguments */
    {
        fprintf(stderr, "Usage: %s <Server IP> <Echo Word> [<Echo Port>]\n", argv[0]);
        exit(1);
    }
}
```

EchoClient.c – parsing the arguments

```
servIP = argv[1];      /* First arg: server IP address (dotted quad) */
echoString = argv[2];  /* Second arg: string to echo */

if ((echoStringLen = strlen(echoString)) > ECHOMAX) /* Check input
length */
    DieWithError("Echo word too long");

if (argc == 4)
    echoServPort = atoi(argv[3]); /* Use given port, if any */
else
    echoServPort = 7; /* 7 is the well-known port for the echo service */
```

EchoClient.c – creating the socket and sending

```
/* Create a datagram/UDP socket */
if ((sock = socket(AF_INET, SOCK_DGRAM, 0)) < 0) DieWithError("socket()
    failed");

/* Construct the server address structure */
memset(&echoServAddr, 0, sizeof(echoServAddr)); /* Zero out structure */
echoServAddr.sin_family = AF_INET; /* Internet addr family */
echoServAddr.sin_addr.s_addr = inet_addr(servIP); /* Server IP address */
echoServAddr.sin_port = htons(echoServPort); /* Server port */

/* Send the string to the server */
if (sendto(sock, echoString, echoStringLen, 0, (struct sockaddr *)
    &echoServAddr, sizeof(echoServAddr)) != echoStringLen)
    DieWithError("sendto() sent a different number of bytes than expected");
```

EchoClient.c – receiving, printing

```
/* Recv a response */
fromSize = sizeof(fromAddr);
if ((respStringLen = recvfrom(sock, echoBuffer, ECHOMAX, 0, (struct
    sockaddr *) &fromAddr, &fromSize)) != echoStringLen)
    DieWithError("recvfrom() failed");
if (echoServAddr.sin_addr.s_addr != fromAddr.sin_addr.s_addr)
    { fprintf(stderr, "Error: received a packet from unknown source.\n");
    exit(1); }

/* null-terminate the received data */
echoBuffer[respStringLen] = '\0';
printf("Received: %s\n", echoBuffer); /* Print the echoed arg */
close(sock);
exit(0);
} /* end of main () */
```

EchoServer.c – #include's and #define's

- `#include <stdio.h>` `/* for printf() and fprintf() */`
- `#include <sys/socket.h>` `/* for socket() and bind() */`
- `#include <arpa/inet.h>` `/* for sockaddr_in and inet_ntoa() */`
- `#include <stdlib.h>` `/* for atoi() and exit() */`
- `#include <string.h>` `/* for memset() */`
- `#include <unistd.h>` `/* for close() */`
- `#define ECHOMAX 255` `/* Longest string to echo */`
- `void DieWithError(char *errorMessage);` `/* External error handling function */`

EchoServer.c – variable declarations and arguments parsing

```
int main(int argc, char *argv[])
{
    int sock;                /* Socket */
    struct sockaddr_in echoServAddr; /* Local address */
    struct sockaddr_in echoCliAddr; /* Client address */
    unsigned int cliAddrLen;    /* Length of incoming message */
    char echoBuffer[ECHOMAX]; /* Buffer for echo string */
    unsigned short echoServPort; /* Server port */
    int recvMsgSize;           /* Size of received message */

    if (argc != 2)            /* Test for correct number of parameters */
    {
        fprintf(stderr, "Usage: %s <UDP SERVER PORT>\n", argv[0]);
        exit(1);
    }

    echoServPort = atoi(argv[1]); /* First arg: local port */
}
```

EchoServer.c – creating and binding socket

```
/* Create socket for sending/receiving datagrams */
if ((sock = socket(AF_INET, SOCK_DGRAM, 0)) < 0)
    DieWithError("socket() failed");

/* Construct local address structure */
memset(&echoServAddr, 0, sizeof(echoServAddr)); /* Zero out
structure */
echoServAddr.sin_family = AF_INET;                /* Internet address
family */
echoServAddr.sin_addr.s_addr = htonl(INADDR_ANY); /* Any
incoming interface */
echoServAddr.sin_port = htons(echoServPort);      /* Local port */

/* Bind to the local address */
if (bind(sock, (struct sockaddr *) &echoServAddr,
sizeof(echoServAddr)) < 0)
    DieWithError("bind() failed");
```


EchoServer.c – do echoing

```
for (;;) /* Run forever */
{
    /* Set the size of the in-out parameter */
    cliAddrLen = sizeof(echoCIntAddr);

    /* Block until receive message from a client */
    if ((recvMsgSize = recvfrom(sock, echoBuffer, ECHOMAX, 0,
        (struct sockaddr *) &echoCIntAddr, &cliAddrLen)) < 0)
        DieWithError("recvfrom() failed");

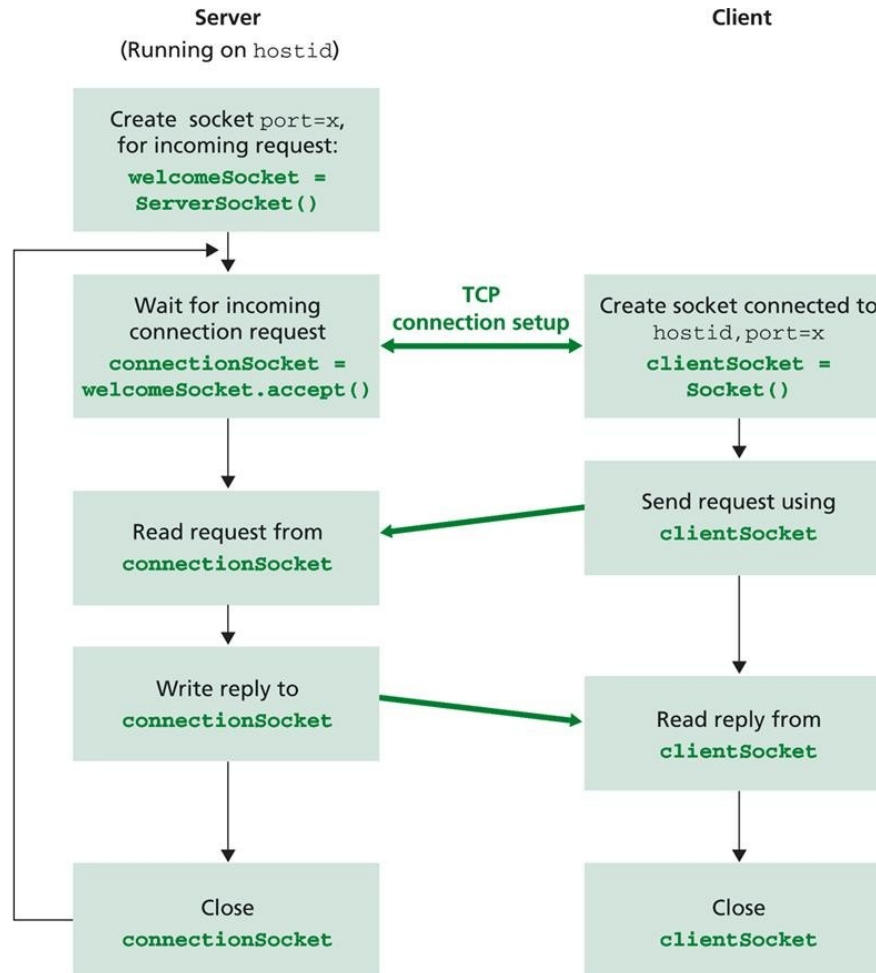
    printf("Handling client %s\n", inet_ntoa(echoCIntAddr.sin_addr));

    /* Send received datagram back to the client */
    if (sendto(sock, echoBuffer, recvMsgSize, 0,
        (struct sockaddr *) &echoCIntAddr, sizeof(echoCIntAddr)) != recvMsgSize)
        DieWithError("sendto() sent a different number of bytes than expected");
}
/* NOT REACHED */
} /* end of main () */
```

An Example Client/Server in Java

- A client reads a line from its **standard input** (keyboard) and sends the line out its socket to the server
- The server reads a line from its connection socket
- The server converts the line to uppercase
- The server sends the modified line out its connection socket to the client
- The client reads the modified line from its socket and prints the line on its **standard output** (monitor)

Main activity between client and server




TCPClient.java

```
import java.io.*;
import java.net.*;
class TCPClient {
    public static void main (String argv[]) throws Exception
    {
        String sentence;
        String modifiedSentence;
        BufferedReader inFromUser = new BufferedReader(new InputStreamReader(System.in));
        Socket clientSocket = new Socket("hostname", 6789); // SOCKET, BIND, CONNECT
        DataOutputStream outToServer = new DataOutputStream(clientSocket.getOutputStream());
        BufferedReader inFromServer = new BufferedReader(new
            InputStreamReader(clientSocket.getInputStream()));
        sentence = inFromUser.readLine();
        outToServer.writeBytes(sentence + '\n');
        modifiedSentence = inFromServer.readLine();
        System.out.println("FROM SERVER: " + modifiedSentence);
        clientSocket.close();
    }
}
```

TCPServer.java

```
import java.io.*;
import java.net.*;
class TCPClient {
    public static void main (String argv[]) throws Exception
    {
        String clientSentence;
        String capitalizedSentence;
        ServerSocket welcomeSocket = new ServerSocket(6789); //SOCKET, BIND, LISTEN
        while(true) {
            Socket connectionSocket = welcomeSocket.accept(); //ACCEPT
            BufferedReader inFromClient = new BufferedReader(new
                InputStreamReader(connectionSocket.getInputStream()));
            DataOutputStream outToClient = new
                DataOutputStream(connectionSocket.getOutputStream());
            clientSentence = inFromClient.readLine(); //RECEIVE
            capitalizedSentence = clientSentence.toUpperCase() + '\n';
            outToClient.writeBytes(capitalizedSentence); //SEND
        }
    }
}
```

Example application: “quote of the day” (Java)



- Client requests the server
- Server sends a quote-of-the-day back to the client
- *We will use datagram sockets*

“Quote of the day”: server

- Creating a socket

```
socket = new DatagramSocket(4445);
```

```
// calls: SOCKET (SOCK_DGRAM), BIND (port 4445) – “we have got a post-address”
```

- Receiving a request from a client

```
byte[] buf = new byte[256]; // new buffer (array of bytes)
```

```
DatagramPacket packet = new DatagramPacket(buf, buf.length); // = a buffer + socket  
specific information - “make a letter box”
```

```
socket.receive(packet); // calls RECEIVE – recvfrom() - “hang the letter box out on the door  
and wait till a letter arrives”
```

- Filling the buffer with another quote (stored in a file)

```
String dString = getNextQuote(); //read another line from a file – “find the requested info”
```

```
buf = dString.getBytes(); // transfer the content of the string into the buffer – “write the letter”
```

- Sending a quote back to the client

```
InetAddress address = packet.getAddress(); // from which IP address did it arrive?
```

```
int port = packet.getPort(); // ... and the port number? – “put the address on the envelop”
```

```
packet = new DatagramPacket(buf, buf.length, address, port); // make up a new packet –  
“pack the letter into the envelop”
```

```
socket.send(packet); //calls: SEND – sendto() - “put the letter to the mail box”
```

“Quote of the day”: client

- Creating a socket

```
socket = new DatagramSocket();  
// calls: SOCKET (SOCK_DGRAM), no BIND
```

- Sending a request to the server

```
byte[] buf = new byte[256];  
InetAddress address = InetAddress.getByName(args[0]); // get the server address from  
the command line  
DatagramPacket packet = new DatagramPacket(buf, buf.length, address, 4445); //  
make up “a letter”  
socket.send(packet); // calls SEND – sendto()
```

- Receiving the respond (a quote line) from the server

```
packet = new DatagramPacket(buf, buf.length); // “make a letter box”  
socket.receive(packet); // calls RECEIVE – recvfrom()  
String received = new String(packet.getData()); // “open the letter” – extract the quote  
System.out.println("Quote of the Moment: " + received); // print the quote
```


Tips for the assignment of this year

with examples in ANSI-C

Tips&Hints

- How would you manage upload and download?

Use multiple threads

- Speed limitations

Do it in the application layer

- Downloading from multiple sources

Both single- and multiple-threaded approaches are possible

General tip

- Always check the returned values (or handle exceptions)

Some links

- UNIX sockets:

[Tutorial](#)

<http://www.ecst.csuchico.edu/~beej/guide/net/html/>

[Examples \(also used in this tutorial\):](#)

<http://cs.baylor.edu/~donahoo/practical/C.Sockets/textcode.html>

- Winsock:

[FAQ](#)

<http://tangentsoft.net/wskfaq/>

Good luck with the assignment

- *More questions?*